

# Data Mining

Practical Machine Learning Tools and Techniques

Slides for Chapter 4, Algorithms: the basic methods

of *Data Mining* by I. H. Witten, E. Frank,  
M. A. Hall and C. J. Pal

# Algorithms: The basic methods

- Inferring rudimentary rules
- Simple probabilistic modeling
- Constructing decision trees
- Constructing rules
- Association rule learning
- Linear models
- Instance-based learning
- Clustering
- Multi-instance learning

# Simplicity first

- Simple algorithms often work very well!
- There are many kinds of simple structure, e.g.:
  - One attribute does all the work
  - All attributes contribute equally & independently
  - Logical structure with a few attributes suitable for tree
  - A set of simple logical rules
  - Relationships between groups of attributes
  - A weighted linear combination of the attributes
  - Strong neighborhood relationships based on distance
  - Clusters of data in unlabeled data
  - Bags of instances that can be aggregated
- Success of method depends on the domain

# Inferring rudimentary rules

- 1R rule learner: learns a 1-level decision tree
  - A set of rules that all test one particular attribute that has been identified as the one that yields the lowest classification error
- Basic version for finding the rule set from a given training set (assumes nominal attributes):
  - For each attribute
    - Make one branch for each value of the attribute
    - To each branch, assign the most frequent class value of the instances pertaining to that branch
    - Error rate: proportion of instances that do not belong to the majority class of their corresponding branch
  - Choose attribute with lowest error rate

# Pseudo-code for 1R

```
For each attribute,  
  For each value of the attribute, make a rule as follows:  
    count how often each class appears  
    find the most frequent class  
    make the rule assign that class to this attribute-value  
Calculate the error rate of the rules  
Choose the rules with the smallest error rate
```

- 1R's handling of missing values: a missing value is treated as a separate attribute value

# Evaluating the weather attributes

| Outlook  | Temp | Humidity | Windy | Play |                   |                |        |              |
|----------|------|----------|-------|------|-------------------|----------------|--------|--------------|
| Sunny    | Hot  | High     | False | No   | Attribute         | Rules          | Errors | Total errors |
| Sunny    | Hot  | High     | True  | No   | Outlook           | Sunny → No     | 2/5    | 4/14         |
| Overcast | Hot  | High     | False | Yes  |                   | Overcast → Yes | 0/4    |              |
| Rainy    | Mild | High     | False | Yes  |                   | Rainy → Yes    | 2/5    |              |
| Rainy    | Cool | Normal   | False | Yes  | Temp              | Hot → No*      | 2/4    | 5/14         |
| Rainy    | Cool | Normal   | True  | No   |                   | Mild → Yes     | 2/6    |              |
| Overcast | Cool | Normal   | True  | Yes  |                   | Cool → Yes     | 1/4    |              |
| Sunny    | Mild | High     | False | No   | Humidity          | High → No      | 3/7    | 4/14         |
| Sunny    | Cool | Normal   | False | Yes  |                   | Normal → Yes   | 1/7    |              |
| Rainy    | Mild | Normal   | False | Yes  | Windy             | False → Yes    | 2/8    | 5/14         |
| Sunny    | Mild | Normal   | True  | Yes  |                   | True → No*     | 3/6    |              |
| Overcast | Mild | High     | True  | Yes  |                   |                |        |              |
| Overcast | Hot  | Normal   | False | Yes  |                   |                |        |              |
| Rainy    | Mild | High     | True  | No   | * indicates a tie |                |        |              |

\* indicates a tie

# Dealing with numeric attributes






- Idea: discretize numeric attributes into sub ranges (intervals)
- How to divide each attribute's overall range into intervals?
  - Sort instances according to attribute's values
  - Place breakpoints where (majority) class changes
  - This minimizes the total classification error
- Example: *temperature* from weather data

64    65    68    69    70    71    72    72    75    75    80    81    83    85  
Yes | No | Yes Yes Yes | No No Yes | Yes Yes | No | Yes Yes | No

| Outlook  | Temperature | Humidity | Windy | Play |
|----------|-------------|----------|-------|------|
| Sunny    | 85          | 85       | False | No   |
| Sunny    | 80          | 90       | True  | No   |
| Overcast | 83          | 86       | False | Yes  |
| Rainy    | 75          | 80       | False | Yes  |
| ...      | ...         | ...      | ...   | ...  |

# The problem of overfitting

- Discretization procedure is very sensitive to noise
  - A single instance with an incorrect class label will probably produce a separate interval
- Also, something like a *time stamp* attribute will have zero errors
- Simple solution:  
*enforce minimum number of instances in majority class per interval*
- Example: *temperature* attribute with required minimum number of instances in majority class set to three:

|     |   |     |   |     |   |    |    |     |     |   |     |     |     |     |   |     |     |    |   |    |
|-----|---|-----|---|-----|---|----|----|-----|-----|---|-----|-----|-----|-----|---|-----|-----|----|---|----|
| 64  | 65  | 68  | 69  | 70  | 71  | 72 | 72 | 75  | 75  | 80  | 81  | 83  | 85  |     |   |     |     |    |   |    |
| Yes |  | No  |  | Yes | Yes   |    | No | No  | Yes |  | Yes | Yes |     | No  |  | Yes | Yes |    |  | No |
| 64  | 65  | 68  | 69  | 70  | 71  | 72 | 72 | 75  | 75  | 80  | 81  | 83  | 85  |     |   |     |     |    |   |    |
| Yes | No  | Yes | Yes   | Yes |  | No | No | Yes | Yes | Yes   |     | No  | Yes | Yes | No  | Yes | Yes | No |   |    |



# Results with overfitting avoidance

- Resulting rule sets for the four attributes in the weather data, with only two rules for the temperature attribute:

| Attribute   | Rules                                   | Errors | Total errors |
|-------------|---|--------|--------------|
| Outlook     | Sunny $\rightarrow$ No                  | 2/5    | 4/14         |
|             | Overcast $\rightarrow$ Yes              | 0/4    |              |
|             | Rainy $\rightarrow$ Yes                 | 2/5    |              |
| Temperature | $\leq 77.5 \rightarrow$ Yes             | 3/10   | 5/14         |
|             | $> 77.5 \rightarrow$ No*                | 2/4    |              |
| Humidity    | $\leq 82.5 \rightarrow$ Yes             | 1/7    | 3/14         |
|             | $> 82.5$ and $\leq 95.5 \rightarrow$ No | 2/6    |              |
|             | $> 95.5 \rightarrow$ Yes                | 0/1    |              |
| Windy       | False $\rightarrow$ Yes                 | 2/8    | 5/14         |
|             | True $\rightarrow$ No*                  | 3/6    |              |

# Discussion of 1R

- 1R was described in a paper by Holte (1993):

**Very Simple Classification Rules Perform Well on Most Commonly Used Datasets**

Robert C. Holte, Computer Science Department, University of Ottawa

- Contains an experimental evaluation on 16 datasets (using *cross-validation* to estimate classification accuracy on fresh data)
  - Required minimum number of instances in majority class was set to 6 after some experimentation
  - 1R's simple rules performed not much worse than much more complex decision trees
- Lesson: simplicity first can pay off on practical datasets
  - Note that 1R does not perform as well on more recent, more sophisticated benchmark datasets

# Simple probabilistic modeling

- “Opposite” of 1R: use all the attributes
- Two assumptions: Attributes are
  - *equally important*
  - *statistically independent* (given the class value)
    - This means knowing the value of one attribute tells us nothing about the value of another takes on (if the class is known)
- Independence assumption is almost never correct!
- But ... this scheme often works surprisingly well in practice
- The scheme is easy to implement in a program and very fast
- It is known as *naïve Bayes*

# Probabilities for weather data

| Outlook       |     |     | Temperature   |     |     | Humidity      |     |     | Windy         |     |     | Play       |           |
|---------------|-----|-----|---------------|-----|-----|---------------|-----|-----|---------------|-----|-----|------------|-----------|
| <i>Yes No</i> |     |     | <i>Yes No</i> |     |     | <i>Yes No</i> |     |     | <i>Yes No</i> |     |     | <i>Yes</i> | <i>No</i> |
| Sunny         | 2   | 3   | Hot           | 2   | 2   | High          | 3   | 4   | False         | 6   | 2   | 9          | 5         |
| Overcast      | 4   | 0   | Mild          | 4   | 2   | Normal        | 6   | 1   | True          | 3   | 3   |            |           |
| Rainy         | 3   | 2   | Cool          | 3   | 1   |               |     |     |               |     |     |            |           |
| Sunny         | 2/9 | 3/5 | Hot           | 2/9 | 2/5 | High          | 3/9 | 4/5 | False         | 6/9 | 2/5 | 9/14       | 5/14      |
| Overcast      | 4/9 | 0/5 | Mild          | 4/9 | 2/5 | Normal        | 6/9 | 1/5 | True          | 3/9 | 3/5 |            |           |
| Rainy         | 3/9 | 2/5 | Cool          | 3/9 | 1/5 |               |     |     |               |     |     |            |           |

| Outlook  | Temp | Humidity | Windy | Play |
|----------|------|----------|-------|------|
| Sunny    | Hot  | High     | False | No   |
| Sunny    | Hot  | High     | True  | No   |
| Overcast | Hot  | High     | False | Yes  |
| Rainy    | Mild | High     | False | Yes  |
| Rainy    | Cool | Normal   | False | Yes  |
| Rainy    | Cool | Normal   | True  | No   |
| Overcast | Cool | Normal   | True  | Yes  |
| Sunny    | Mild | High     | False | No   |
| Sunny    | Cool | Normal   | False | Yes  |
| Rainy    | Mild | Normal   | False | Yes  |
| Sunny    | Mild | Normal   | True  | Yes  |
| Overcast | Mild | High     | True  | Yes  |
| Overcast | Hot  | Normal   | False | Yes  |
| Rainy    | Mild | High     | True  | No   |

# Probabilities for weather data

| Outlook  |     |     | Temperature |     |     | Humidity |     |     | Windy |     |     | Play |      |
|----------|-----|-----|-------------|-----|-----|----------|-----|-----|-------|-----|-----|------|------|
|          | Yes | No  |             | Yes | No  |          | Yes | No  |       | Yes | No  | Yes  | No   |
| Sunny    | 2   | 3   | Hot         | 2   | 2   | High     | 3   | 4   | False | 6   | 2   | 9    | 5    |
| Overcast | 4   | 0   | Mild        | 4   | 2   | Normal   | 6   | 1   | True  | 3   | 3   |      |      |
| Rainy    | 3   | 2   | Cool        | 3   | 1   |          |     |     |       |     |     |      |      |
| Sunny    | 2/9 | 3/5 | Hot         | 2/9 | 2/5 | High     | 3/9 | 4/5 | False | 6/9 | 2/5 | 9/14 | 5/14 |
| Overcast | 4/9 | 0/5 | Mild        | 4/9 | 2/5 | Normal   | 6/9 | 1/5 | True  | 3/9 | 3/5 |      |      |
| Rainy    | 3/9 | 2/5 | Cool        | 3/9 | 1/5 |          |     |     |       |     |     |      |      |

- A new day:

| Outlook | Temp. | Humidity | Windy | Play |
|---------|-------|----------|-------|------|
| Sunny   | Cool  | High     | True  | ?    |

Likelihood of the two classes

$$\text{For "yes"} = 2/9 \times 3/9 \times 3/9 \times 3/9 \times 9/14 = 0.0053$$

$$\text{For "no"} = 3/5 \times 1/5 \times 4/5 \times 3/5 \times 5/14 = 0.0206$$

Conversion into a probability by normalization:

$$P(\text{"yes"}) = 0.0053 / (0.0053 + 0.0206) = 0.205$$

$$P(\text{"no"}) = 0.0206 / (0.0053 + 0.0206) = 0.795$$

# Can combine probabilities using Bayes's rule

- Famous rule from probability theory due to

**Thomas Bayes**

**Born: 1702 in London, England**

**Died: 1761 in Tunbridge Wells, Kent, England**

- Probability of an event  $H$  given observed evidence  $E$ :

$$P(H | E) = P(E | H)P(H) / P(E)$$

- *A priori* probability of  $H$  :  $P(H)$ 
  - Probability of event *before* evidence is seen
- *A posteriori* probability of  $H$  :  $P(H | E)$ 
  - Probability of event *after* evidence is seen

# Naïve Bayes for classification

- Classification learning: what is the probability of the class given an instance?
  - Evidence  $E$  = instance's non-class attribute values
  - Event  $H$  = class value of instance
- Naïve assumption: evidence splits into parts (i.e., attributes) that are conditionally *independent*
- This means, given  $n$  attributes, we can write Bayes' rule using a product of per-attribute probabilities:

$$P(H | E) = P(E_1 | H)P(E_2 | H) \cdots P(E_n | H)P(H) / P(E)$$

# Weather data example

| Outlook | Temp. | Humidity | Windy | Play |
|---------|-------|----------|-------|------|
| Sunny   | Cool  | High     | True  | ?    |

← *Evidence E*

*Probability of class “yes”* →

$$\begin{aligned} P(\text{yes} \mid E) &= P(\text{Outlook} = \text{Sunny} \mid \text{yes}) \\ &\quad P(\text{Temperature} = \text{Cool} \mid \text{yes}) \\ &\quad P(\text{Humidity} = \text{High} \mid \text{yes}) \\ &\quad P(\text{Windy} = \text{True} \mid \text{yes}) \\ &\quad P(\text{yes}) / P(E) \\ &= \frac{2/9 \times 3/9 \times 3/9 \times 3/9 \times 9/14}{P(E)} \end{aligned}$$



# The “zero-frequency problem”

- What if an attribute value does not occur with every class value?  
(e.g., “Humidity = high” for class “yes”)
  - Probability will be zero:  $P(\text{Humidity} = \text{High} \mid \text{yes}) = 0$
  - *A posteriori* probability will also be zero:  $P(\text{yes} \mid E) = 0$   
(Regardless of how likely the other values are!)
- Remedy: add 1 to the count for every attribute value-class combination (*Laplace estimator*)
- Result: probabilities will never be zero
- Additional advantage: stabilizes probability estimates computed from small samples of data

# Modified probability estimates

- In some cases adding a constant different from 1 might be more appropriate
- Example: attribute *outlook* for class *yes*

$$\frac{2 + \mu/3}{9 + \mu}$$

*Sunny*

$$\frac{4 + \mu/3}{9 + \mu}$$

*Overcast*

$$\frac{3 + \mu/3}{9 + \mu}$$

*Rainy*

- Weights don't need to be equal (but they must sum to 1)

$$\frac{2 + \mu p_1}{9 + \mu}$$

$$\frac{4 + \mu p_2}{9 + \mu}$$

$$\frac{3 + \mu p_3}{9 + \mu}$$

# Missing values

- Training: instance is not included in frequency count for attribute value-class combination
- Classification: attribute will be omitted from calculation
- Example:

| Outlook | Temp. | Humidity | Windy | Play |
|---------|-------|----------|-------|------|
| ?       | Cool  | High     | True  | ?    |

Likelihood of "yes" =  $3/9 \times 3/9 \times 3/9 \times 9/14 = 0.0238$

Likelihood of "no" =  $1/5 \times 4/5 \times 3/5 \times 5/14 = 0.0343$

$P(\text{"yes"}) = 0.0238 / (0.0238 + 0.0343) = 41\%$

$P(\text{"no"}) = 0.0343 / (0.0238 + 0.0343) = 59\%$

# Numeric attributes

- Usual assumption: attributes have a *normal* or *Gaussian* probability distribution (given the class)
- The *probability density function* for the normal distribution is defined by two parameters:


- *Sample mean*  $\mu$

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i$$

- *Standard deviation*  $\sigma$

$$\sigma = \sqrt{\frac{1}{N-1} \sum_{i=1}^N (x_i - \mu)^2}$$

- Then the density function  $f(x)$  is  $f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$



# Statistics for weather data

| Outlook       |     |     | Temperature    |                | Humidity        |                | Windy         |     |     | Play          |    |
|---------------|-----|-----|----------------|----------------|-----------------|----------------|---------------|-----|-----|---------------|----|
| <i>Yes No</i> |     |     | <i>Yes No</i>  |                | <i>Yes No</i>   |                | <i>Yes No</i> |     |     | <i>Yes No</i> |    |
| Sunny         | 2   | 3   | 64, 68,        | 65,71,         | 65, 70,         | 70, 85,        | False         | 6   | 2   | 9             | 5  |
| Overcast      | 4   | 0   | 69, 70,        | 72,80,         | 70, 75,         | 90, 91,        | True          | 3   | 3   |               |    |
| Rainy         | 3   | 2   | 72, ...        | 85, ...        | 80, ...         | 95, ...        |               |     |     |               |    |
| Sunny         | 2/9 | 3/5 | $\mu = 73$     | $\mu = 75$     | $\mu = 79$      | $\mu = 86$     | False         | 6/9 | 2/5 | 9/            | 5/ |
| Overcast      | 4/9 | 0/5 | $\sigma = 6.2$ | $\sigma = 7.9$ | $\sigma = 10.2$ | $\sigma = 9.7$ | True          | 3/9 | 3/5 | 14            | 14 |
| Rainy         | 3/9 | 2/5 |                |                |                 |                |               |     |     |               |    |

- Example density value:

$$f(\text{temperature} = 66|\text{yes}) = \frac{1}{\sqrt{2\pi} \cdot 6.2} e^{-\frac{(66-73)^2}{2 \cdot 6.2^2}} = 0.0340$$

# Classifying a new day

- A new day:

| Outlook | Temp. | Humidity | Windy | Play |
|---------|-------|----------|-------|------|
| Sunny   | 66    | 90       | true  | ?    |

Likelihood of "yes" =  $2/9 \times 0.0340 \times 0.0221 \times 3/9 \times 9/14 = 0.000036$

Likelihood of "no" =  $3/5 \times 0.0221 \times 0.0381 \times 3/5 \times 5/14 = 0.000108$

$P(\text{"yes"}) = 0.000036 / (0.000036 + 0.000108) = 25\%$

$P(\text{"no"}) = 0.000108 / (0.000036 + 0.000108) = 75\%$

- Missing values during training are not included in calculation of mean and standard deviation

# Probability densities

- Probability densities  $f(x)$  can be greater than 1; hence, they are not probabilities
  - However, they must integrate to 1: the area under the probability density curve must be 1
- Approximate relationship between probability and probability density can be stated as

$$P(x - \varepsilon / 2 \leq X \leq x + \varepsilon / 2) \approx \varepsilon f(x)$$

assuming  $\varepsilon$  is sufficiently small

- When computing likelihoods, we can treat densities just like probabilities

# Naïve Bayes: discussion


- Naïve Bayes works surprisingly well even if independence assumption is clearly violated
- Why? Because classification does not require accurate probability estimates *as long as maximum probability is assigned to the correct class*
- However: adding too many redundant attributes will cause problems (e.g., identical attributes)
- Note also: many numeric attributes are not normally distributed (*kernel density estimators* can be used instead)



# Covering algorithms

- Can convert decision tree into a rule set
  - Straightforward, but rule set overly complex
  - More effective conversions are not trivial and may incur a lot of computation
- Instead, we can generate rule set directly
  - One approach: for each class in turn, find rule set that covers all instances in it  
(excluding instances not in the class)
- Called a *covering* approach:
  - At each stage of the algorithm, a rule is identified that “covers” some of the instances

# Example: generating a rule



- Possible rule set for class “b”:


`If  $x \leq 1.2$  then class = b`

`If  $x > 1.2$  and  $y \leq 2.6$  then class = b`

- Could add more rules, get “perfect” rule set

# Rules vs. trees


- Corresponding decision tree:  
(produces exactly the same predictions)



- But: rule sets *can* be more perspicuous when decision trees suffer from replicated subtrees
- Also: in multiclass situations, covering algorithm concentrates on one class at a time whereas decision tree learner takes all classes into account

# Simple covering algorithm

- Basic idea: generate a rule by adding tests that maximize the rule's accuracy
- Similar to situation in decision trees: problem of selecting an attribute to split on
  - But: decision tree inducer maximizes overall purity
- Each new test reduces rule's coverage:



# Selecting a test

- Goal: maximize accuracy
  - $t$  total number of instances covered by rule
  - $p$  positive examples of the class covered by rule
  - $t - p$  number of errors made by rule
  - Select test that maximizes the ratio  $p/t$
- We are finished when  $p/t = 1$  or the set of instances cannot be split any further

# Example: contact lens data

- Rule we seek:

```
If ?  
    then recommendation = hard
```

- Possible tests:

|                                       |      |
|---------------------------------------|------|
| Age = Young                           | 2/8  |
| Age = Pre-presbyopic                  | 1/8  |
| Age = Presbyopic                      | 1/8  |
| Spectacle prescription = Myope        | 3/12 |
| Spectacle prescription = Hypermetrope | 1/12 |
| Astigmatism = no                      | 0/12 |
| Astigmatism = yes                     | 4/12 |
| Tear production rate = Reduced        | 0/12 |
| Tear production rate = Normal         | 4/12 |

# Modified rule and resulting data

- Rule with best test added:

```
If astigmatism = yes  
    then recommendation = hard
```

- Instances covered by modified rule:

| Age            | Spectacle prescription | Astigmatism | Tear production rate | Recommended lenses |
|----------------|------------------------|-------------|----------------------|--------------------|
| Young          | Myope                  | Yes         | Reduced              | None               |
| Young          | Myope                  | Yes         | Normal               | Hard               |
| Young          | Hypermetrope           | Yes         | Reduced              | None               |
| Young          | Hypermetrope           | Yes         | Normal               | hard               |
| Pre-presbyopic | Myope                  | Yes         | Reduced              | None               |
| Pre-presbyopic | Myope                  | Yes         | Normal               | Hard               |
| Pre-presbyopic | Hypermetrope           | Yes         | Reduced              | None               |
| Pre-presbyopic | Hypermetrope           | Yes         | Normal               | None               |
| Presbyopic     | Myope                  | Yes         | Reduced              | None               |
| Presbyopic     | Myope                  | Yes         | Normal               | Hard               |
| Presbyopic     | Hypermetrope           | Yes         | Reduced              | None               |
| Presbyopic     | Hypermetrope           | Yes         | Normal               | None               |

# Further refinement

- Current state:

```
If astigmatism = yes  
    and ?  
    then recommendation = hard
```

- Possible tests:

|                                       |       |
|---------------------------------------|-------|
| Age = Young                           | 2 / 4 |
| Age = Pre-presbyopic                  | 1 / 4 |
| Age = Presbyopic                      | 1 / 4 |
| Spectacle prescription = Myope        | 3 / 6 |
| Spectacle prescription = Hypermetrope | 1 / 6 |
| Tear production rate = Reduced        | 0 / 6 |
| Tear production rate = Normal         | 4 / 6 |



# Modified rule and resulting data

- Rule with best test added:

```
If astigmatism = yes
    and tear production rate = normal
then recommendation = hard
```

- Instances covered by modified rule:

| Age            | Spectacle prescription | Astigmatism | Tear production rate | Recommended lenses |
|----------------|------------------------|-------------|----------------------|--------------------|
| Young          | Myope                  | Yes         | Normal               | Hard               |
| Young          | Hypermetrope           | Yes         | Normal               | hard               |
| Pre-presbyopic | Myope                  | Yes         | Normal               | Hard               |
| Pre-presbyopic | Hypermetrope           | Yes         | Normal               | None               |
| Presbyopic     | Myope                  | Yes         | Normal               | Hard               |
| Presbyopic     | Hypermetrope           | Yes         | Normal               | None               |

# Further refinement

- Current state:

```
If astigmatism = yes
    and tear production rate = normal
    and ?
    then recommendation = hard
```

- Possible tests:

|                                       |     |
|---------------------------------------|-----|
| Age = Young                           | 2/2 |
| Age = Pre-presbyopic                  | 1/2 |
| Age = Presbyopic                      | 1/2 |
| Spectacle prescription = Myope        | 3/3 |
| Spectacle prescription = Hypermetrope | 1/3 |

- Tie between the first and the fourth test
  - We choose the one with greater coverage

# The final rule

- Final rule:

```
If astigmatism = yes  
and tear production rate = normal  
and spectacle prescription = myope  
then recommendation = hard
```


- Second rule for recommending “hard lenses”:  
(built from instances not covered by first rule)

```
If age = young and astigmatism = yes  
and tear production rate = normal  
then recommendation = hard
```

- These two rules cover all “hard lenses”:
  - Process is repeated with other two classes

# Pseudo-code for PRISM

```
For each class C
  Initialize E to the instance set
  While E contains instances in class C
    Create a rule R with an empty left-hand side that predicts class C
    Until R is perfect (or there are no more attributes to use) do
      For each attribute A not mentioned in R, and each value v,
        Consider adding the condition A = v to the left-hand side of R
        Select A and v to maximize the accuracy p/t
        (break ties by choosing the condition with the largest p)
      Add A = v to R
    Remove the instances covered by R from E
```



# Rules vs. decision lists

- PRISM with outer loop removed generates a decision list for one class
  - Subsequent rules are designed for rules that are not covered by previous rules
  - But: order does not matter because all rules predict the same class so outcome does not change if rules are shuffled
- Outer loop considers all classes separately
  - No order dependence implied
- Problems: overlapping rules, default rule required

# Separate and conquer rule learning

- Rule learning methods like the one PRISM employs (for each class) are called *separate-and-conquer* algorithms:
  - First, identify a useful rule
  - Then, separate out all the instances it covers
  - Finally, “conquer” the remaining instances
- Difference to divide-and-conquer methods:
  - Subset covered by a rule does not need to be explored any further